

V.Ger Travel Company - An Analysis

ITNBD4 Assignment

Student: 2710017

13 January 2025

Table of Contents

1	Summary	3
1.1	Approach & Findings	3
1.2	Next Steps	3
2	Introduction	4
2.1	Background & Approach	4
2.2	Results & Further Details	4
3	Hotel Demand Forecasting	5
3.1	Business Scope & Benefits	5
3.2	Data Analysis Approach	6
3.3	Simulated Data Analysis	8
3.4	Conclusions & Next Steps	14
4	Customer Satisfaction	15
4.1	Business Scope & Benefits	15
4.2	Data Analysis Approach	15
4.3	Simulated Data Analysis	17
4.4	Conclusions & Next Steps	21
	References	22
	Appendices	24
A	Hotel Demand Forecasting - Jupyter Notebook Output	24
A.1	Data Load & Characteristics	24
A.2	Ordinary Least Squares (OLS) Linear Regression	28
A.3	SARIMA Model Creation & Forecasting	30
B	SARIMA Models Creation + Auto ARIMA - Jupyter Notebook Output	33
B.1	SARIMA Model Create & Evaluate - Manual	33
B.2	SARIMA Model Create & Evaluate - Auto ARIMA - First	35
B.3	SARIMA Model Create & Evaluate - Auto ARIMA - Second	40
C	Customer Satisfaction - Jupyter Notebook Output	44
C.1	Data Load & Characteristics	44
C.2	XGBoost Model Creation & Forecasting	45

1 Summary

1.1 Approach & Findings

Summary of the report, sort of abstract type summary?

1.2 Next Steps

Summary of conclusions and next steps

2 Introduction

2.1 Background & Approach

The travel conglomerate V.Ger Travel (VGT) has a broad range of operations, including hotels, resorts, car rentals and also air travel through charter flights. Travel bookings originate primarily from VGT's own travel web site which is supported by several operational information systems that cover all aspects of it's business from customer relations through to logistics and maintenance. There is a wealth of data available with over 10 years history of travel bookings, however, the use of modern Data Science methods to harness this data is in its infancy at VGT.

This report describes the recommendations of the new Chief Data Officer (CDO) as to how to implement modern Data Science techniques to utilise VGT's data to improve its efficiency and profit. Having completed an initial high level review, two areas were identified that appeared to offer the biggest opportunities for increasing efficiency and also that are achievable as a first step in implementing new techniques. The two use cases are:

- Hotel Demand Forecasting
- Customer Satisfaction

Each of these use cases are explored in the following sections, each of which describes:

- Business Scope & Benefits - What is proposed and why is this helpful
- Data Analysis Approach - The suggested analysis techniques and the data required
- Simulated Data Analysis - An example of the analysis using simulated data
- Conclusions & Next Steps - Findings and how to move forward

2.2 Results & Further Details

A summary of the approach taken and key results of the analysis are described for each use case in the individual sections of this report. Additionally, all details for generating the simulated data and the analysis steps are described at the following:

- Appendices - Step by step approach, with data tables and plots
- Jupyter Notebooks - Full listings of Python code and results in a [GitHub Repository](#)
- Data Files - Generated CSV files in a [GitHub Repository](#)

3 Hotel Demand Forecasting

3.1 Business Scope & Benefits

The hotel operations of VGT are a significant part of its business and any efficiencies in this area have the potential to make large contributions to overall profitability. The profit contribution from individual hotels can be maximised by ensuring revenue is as high as possible whilst at the same time minimising a hotel's operating costs. One way to do this is to provide a hotel's management team with the tools to carry out reliable demand forecasting.

Hotel demand forecasting is the prediction of the demand for rooms and related hotel services to help a hotel's management team determine pricing, staffing and marketing strategies (Johansson, 2022). If the demand for hotel rooms can be reliably forecast then this enables:

- **Dynamic Pricing** - Adjust future prices in response to forecast demand. When high demand is expected then future room rates can be increased; when low demand is expected then discounts can be offered or packages can be advertised. And marketing strategies can be determined to respond to the demand forecasts
- **Staffing Levels** - Hotel staffing can be adjusted to maintain customer service levels but not over staff when demand is expected to be lower
- **Inventory Management** - Similarly inventory can be adjusted, for example catering supplies maintained just sufficient to meet the forecast number of customers using catering facilities.

There are many factors that will influence the demand for hotel rooms, some may remain stable whilst others are less so and will vary over time or in response to external factors, these include:

- **Location & Market** - Is the hotel a budget or a boutique hotel? Is it in a business district, near a beach or a ski resort etc. Are customers mainly business travellers or tourists?
- **Economic** - Macro level impact from the state of the economy
- **Local Competition** - Competition with local hotels
- **Seasonality** - When are the high and low seasons? Is it a summer or a winter resort? What are the local weather patterns and seasons. Are there weather related attractions?
- **Local Events** - When are any local festivities, sports events, school holidays, religious events, music conferences, business conferences?

3.1.1 Occupancy & Revenue Indicators

An important indicator of demand is the occupancy rate (Jeffrey and Hubbard, 1994) and (FHA, 2023) which is simply the percentage of the total rooms occupied for a given time period. The occupancy rate varies across the industry but a target of 60% to 80% is typical.

Occupancy is used alongside revenue related indicators to provide a measure of revenue health, the three main indicators (for a given time period, eg daily) are:

- Occupancy Rate (OCC%) - Percentage of available rooms that are occupied, or expected to be occupied
- Average Daily Rate (ADR) - Average revenue per room occupied, across all room price bands
- Revenue Per Available Room (RevPAR) - Revenue reflecting all available rooms. Calculated by: $OCC\% * ADR$. A good overall indicator of revenue.

3.1.2 Forecasting Occupancy & Business Benefits

The time period used for forecasting will vary depending on the objectives desired (Lighthouse, 2024) and (Lighthouse, 2023) for example:

- Short Term - Forecast occupancy for next month so room pricing can be adjusted appropriately.
- Long-term - Forecast occupancy for next year so price bands and packages can be set, marketing strategies defined and required staffing levels determined.

To recap, if occupancy (and the associated revenue indicators) can be reliably forecast then plans can be put in place to maximise revenue by adjusting pricing and marketing strategies and controlling costs by flexing staffing and inventory levels.

3.2 Data Analysis Approach

At VGT, no rigorous forecasting is currently in place so to begin with a relatively simple approach will be implemented in a small number of hotels. If the benefits of this are confirmed, then it can then be extended in sophistication using more complex forecasting models and for longer time periods. It can then be implemented across all hotels in VGT.

The objective of this first step is to establish a model that can be used to forecast the daily OCC% at an individual hotel for the coming month, ie the forecast is for 30 to 40 days in the future. The forecast will then be used by hotel's management team to: i) adopt dynamic pricing; ii) execute supporting short-term marketing; iii) fine-tune staffing rotas and holiday leave for the coming weeks. This should improve the efficiency of the hotel's operations by increasing room bookings whilst ensuring staffing costs are controlled at an appropriate level.

3.2.1 Forecasting Model & Data Required

The scope of the envisaged forecasting model is to calculate a month of daily OCC% for each room category in an hotel. The output of the forecast model will be provided in spreadsheet form so that the hotel management team can manually make adjustments to try to improve revenue and staffing levels. The actual occupancy and revenue can then be tracked against the forecast throughout the month in order to assess how accurate the forecasting model is and to help refine it.

Forecasting Model Data

The data required for the forecasting model is the last 4 years of daily room occupancy. In this analysis a single hotel with two classes of room (standard and premium) will be used. A 4 year history of daily room bookings was selected with 3 years used for training and 1 year for

validation. This length of history was chosen as a starting point because a previous occupancy study (Phumchusri and Suwatanapongched, 2023) found that the choice could be quite dependent on the scenario; however, there needs to be a balance of too short and missing seasonality vs too long and not being sufficiently responsive. Phumchusri and Suwatanapongched (2023) also found that using 4 years history of daily occupancy to forecast 2 to 8 weeks was a good approach.

The specific data required is:

- For an individual hotel and each room category (standard and premium)
- Daily
- Room capacity
- Room rate
- Rooms occupied

Tracking Spreadsheets

The data comprising the revenue forecasting and tracking spreadsheet is:

- Forecast OCC% (derived from the forecasting model)
- Daily room rates by room category that can be manually adjusted
- Daily Revenue, ADR, RevPAR (derived from the OCC% and room rates)
- Special events, a facility to mark local events that may impact demand. For example, sports events, concerts, conference, unusual weather forecasts
- Actual OCC%, Revenue, ADR, RevPAR

The data provided for the staffing forecasting spreadsheet is:

- Forecast OCC% (from the forecasting model)
- Averaged staff requirements per room
- Staffing levels (derived from the OCC%)
- Actual staffing levels

3.2.2 Techniques Considered

There are several potential tools that can be used with historical time series data to forecast future occupancy rates, ranging from established ‘simpler’ techniques such as linear regression through to more sophisticated machine learning and neural network models (Huang and Zheng, 2022). However, given the relative infancy of Data Science at VGT, the use of more sophisticated tools will be prioritised for future work. In this investigation, the following techniques will be examined:

1. Ordinary Least Squares (OLS) linear regression
2. ARIMA, SARIMA - Use to fully incorporate seasonality
3. SARIMAX - If exogenous factors need to be accounted for

OLS Linear Regression

Investigating techniques for reliably identifying the factors influencing hotel occupancy has been ongoing for several years. See Andrew, Cranage and Lee (1990) and Jeffrey and Hubbard (1994) for earlier work focusing on the use of regression analysis of time series data. Although it is likely that the nature of hotel business means that occupancy will be seasonal, linear regression will still be investigated first.

ARIMA, SARIMA

Occupancy forecasting using ARIMA using factors such as room capacity and marketing expenditure has successfully been used (Chow, Shyu and Wang, 1998). A comparison of forecasting methods (Weatherford and Kimes, 2003) included historical time series analysis of occupancy using ARIMA. The best time period used was not clear and it is a balance of too short and missing seasonality vs too long and not being sufficiently responsive. Also the best analysis method appears to depend on the characteristics of individual hotels and hotel chains. Using a SARIMA approach using 4 years history of daily occupancy to forecast 2 to 8 weeks was found to be a strong approach (Phumchusri and Suwatanapongched, 2023).

SARIMAX

Due to limitations in processing capacity for daily data a SARIMAX model was not created. But see Section 3.4 for next steps.

3.3 Simulated Data Analysis

3.3.1 Data Summary

A dataset was created to simulate the data as defined in the earlier section. The data elements are shown in the tables below and in more detail at Appendix A.

Table 3.1: Loaded File Data Types

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
Standard_OCC	1461	0	0	177	int64	0	1461	0	0
Standard_Capacity	1461	0	0	1	int64	0	1461	0	0
Standard_Rate	1461	0	0	1	int64	0	1461	0	0
Premium_OCC	1461	0	0	101	int64	0	1461	0	0
Premium_Capacity	1461	0	0	1	int64	0	1461	0	0
Premium_Rate	1461	0	0	1	int64	0	1461	0	0

Table 3.2: Loaded File Characteristics

	count	mean	std	min	25%	50%	75%	max
Standard_OCC	1461.00	135.65	50.22	38.00	88.00	136.00	183.00	254.00

	count	mean	std	min	25%	50%	75%	max
Standard_Capacity	1461.00	254.00	0.00	254.00	254.00	254.00	254.00	254.00
Standard_Rate	1461.00	325.00	0.00	325.00	325.00	325.00	325.00	325.00
Premium_OCC	1461.00	56.14	29.35	0.00	31.00	57.00	83.00	100.00
Premium_Capacity	1461.00	100.00	0.00	100.00	100.00	100.00	100.00	100.00
Premium_Rate	1461.00	575.00	0.00	575.00	575.00	575.00	575.00	575.00

3.3.2 Time Series Characteristics

The occupancy time series for the two categories of room are shown in the figure below. This indicates that the occupancy has a strong seasonality with an annual peak and trough; there are also regular spikes to full capacity bookings. The premium room occupancy hits maximum and zero occupancy on several occasions.




Figure 3.1: Occupancy for Each Room Category

The two occupancy time series were examined further (see detail at [Appendix A](#)) using lag plots, ACF plots and ADF tests which indicated autocorrelation, annual seasonality and a positive trend. Differencing was also completed to confirm non-stationarity. Finally a decomposition was completed and this confirmed the seasonality and trend, see the figures below for standard and premium rooms.




Figure 3.2: Occupancy Time Series Decomposition

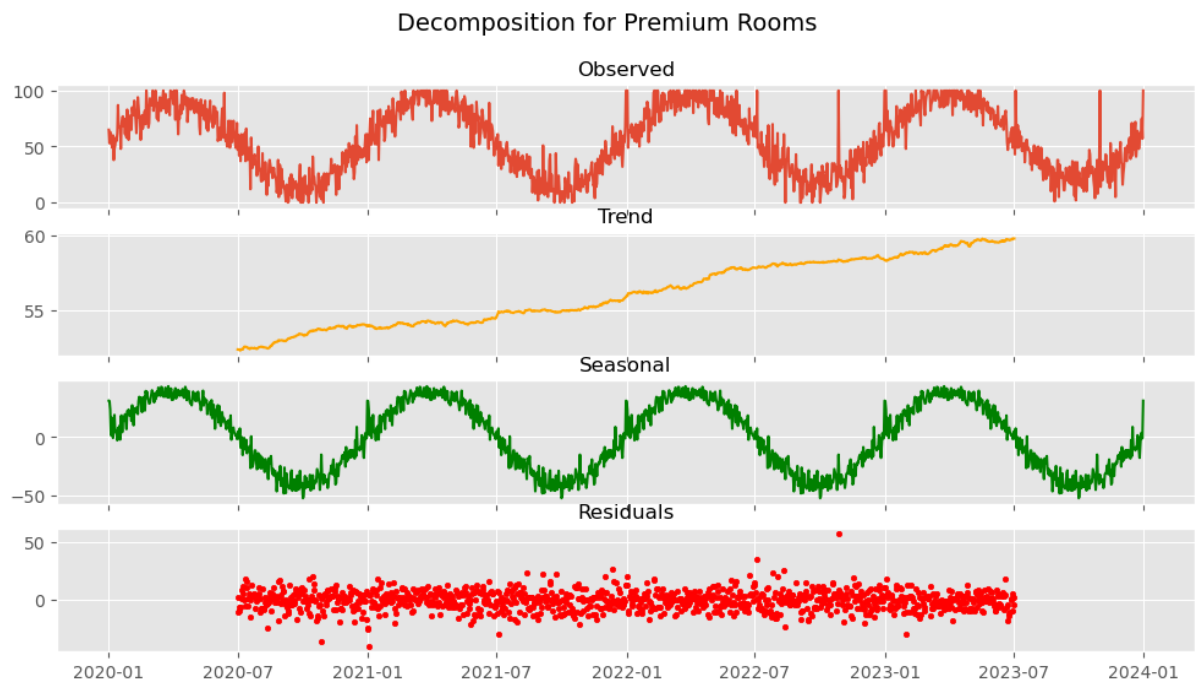


Figure 3.3: Occupancy Time Series Decomposition

3.3.3 OLS Linear Regression

Given the strong seasonality, linear regression is unlikely to be a good model for forecasting, however a regression fit was calculated to double-check. See the plot below which shows that the fit lines could not reliably provide a forecast and even show a downward trend. And the Durban Watson statistics are less than 2 which confirms evidence of autocorrelation.




Figure 3.4: OLS Linear Regression

3.3.4 ARIMA, SARIMA

Given the clear strong seasonality, an ARIMA model was attempted. Parameters for SARIMA were first manually estimated using ACF and PACF plots see figure below and further detail at Appendix A. This suggested the parameters for SARIMA(p,d,q)(PDW)m should have the values of SARIMA(3,60,1)(3,1,1)365. With m-365 reflecting the daily time-steps having an annual seasonal cycle.

An attempt was made to generate a SARIMA model with these parameters but it could not be completed due to insufficient processing capacity. So instead (for the purposes of completing the process) a model was generated with SARIMA(3,1,1)(3,1,1)7. The AIC value for this was 8454.8 and the residuals looked good, see details at Appendix B. However the forecasts produced using the 1 year test data were not accurate, see the figure below, with a MAPE of 259.

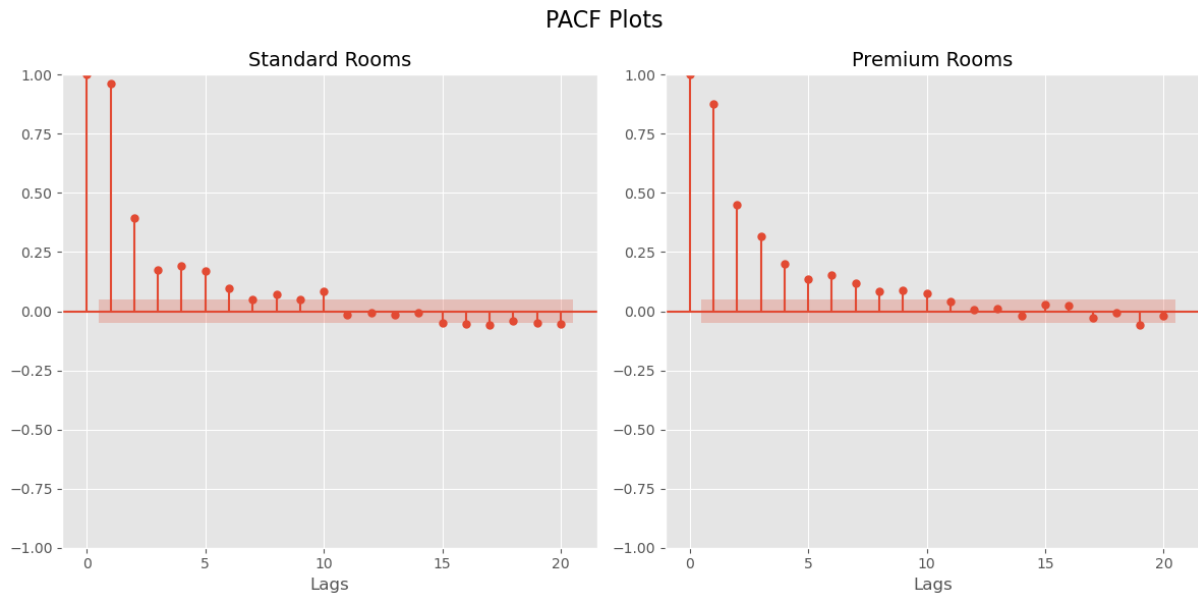


Figure 3.5: Partial Autocorrelation Function Plot




Figure 3.6: Test Data - Forecast vs Actual

Auto ARIMA was also used to attempt to find a better SARIMA model. Again, a seasonal cycle of 365 could not be used due to lack of processing capacity and so $m=7$ was used again. The Auto ARIMA results gave $\text{SARIMA}(4,1,2)(1,1,2)_7$. The AIC value for this was 8444.93 and the residuals looked good, see details at Appendix B. However the forecasts using the 1 year test data were slightly worse, with a MAPE of 308.

3.3.5 Occupancy Forecast

Using the model SARIMA(4,1,2)(1,1,2)7 a one month forecast was produced, see figure below. However, this will not be a reliable model for forecasting given the inaccuracies identified earlier.




Figure 3.7: One Month Occupancy Forecast

If a reliable one month forecast could be created then the OCC% data would then be used in an occupancy spreadsheet which the hotel management team could use for demand pricing for the coming month, as in the example below.

Date	Rooms - Standard				Rooms - Premium				Hotel			
	Capacity	Rate £	Occupied	OCC%	Capacity	Rate £	Occupied	OCC%	OCC%	Revenue	ADR	RevPAR
01/01/2025	30	£150	10	33.3%	10	£225	5	50.0%	37.5%	£2,625	£175	£66
02/01/2025	30	£150	15	50.0%	10	£225	5	50.0%	50.0%	£3,375	£169	£84
03/01/2025	30	£150	21	70.0%	10	£225	2	20.0%	57.5%	£3,600	£157	£90
04/01/2025	30	£150	5	16.7%	10	£225	0	0.0%	12.5%	£750	£150	£19
05/01/2025	30	£150	30	100.0%	10	£225	3	30.0%	82.5%	£5,175	£157	£129
06/01/2025	30	£150	30	100.0%	10	£225	3	30.0%	82.5%	£5,175	£157	£129
07/01/2025	30	£150	15	50.0%	10	£225	8	80.0%	57.5%	£4,050	£176	£101
08/01/2025	30	£150	10	33.3%	10	£225	9	90.0%	47.5%	£3,525	£186	£88
09/01/2025	30	£150	15	50.0%	10	£225	10	100.0%	62.5%	£4,500	£180	£113
10/01/2025	30	£150	21	70.0%	10	£225	2	20.0%	57.5%	£3,600	£157	£90
11/01/2025	30	£150	5	16.7%	10	£225	4	40.0%	22.5%	£1,650	£183	£41
12/01/2025	30	£150	30	100.0%	10	£225	10	100.0%	100.0%	£6,750	£169	£169
13/01/2025	30	£150	30	100.0%	10	£225	10	100.0%	100.0%	£6,750	£169	£169
14/01/2025	30	£150	15	50.0%	10	£225	2	20.0%	42.5%	£2,700	£159	£68
15/01/2025	30	£150	10	33.3%	10	£225	0	0.0%	25.0%	£1,500	£150	£38
16/01/2025	30	£150	15	50.0%	10	£225	3	30.0%	45.0%	£2,925	£163	£73
17/01/2025	30	£150	21	70.0%	10	£225	3	30.0%	60.0%	£3,825	£159	£96
18/01/2025	30	£150	5	16.7%	10	£225	8	80.0%	32.5%	£2,550	£196	£64
19/01/2025	30	£150	30	100.0%	10	£225	9	90.0%	97.5%	£6,525	£167	£163
20/01/2025	30	£150	30	100.0%	10	£225	10	100.0%	100.0%	£6,750	£169	£169
21/01/2025	30	£150	15	50.0%	10	£225	2	20.0%	42.5%	£2,700	£159	£68
22/01/2025	30	£150	10	33.3%	10	£225	4	40.0%	35.0%	£2,400	£171	£60
23/01/2025	30	£150	15	50.0%	10	£225	5	50.0%	50.0%	£3,375	£169	£84
24/01/2025	30	£150	21	70.0%	10	£225	5	50.0%	65.0%	£4,275	£164	£107
25/01/2025	30	£150	5	16.7%	10	£225	2	20.0%	17.5%	£1,200	£171	£30
26/01/2025	30	£150	30	100.0%	10	£225	0	0.0%	75.0%	£4,500	£150	£113
27/01/2025	30	£150	30	100.0%	10	£225	3	30.0%	82.5%	£5,175	£157	£129
28/01/2025	30	£150	15	50.0%	10	£225	3	30.0%	45.0%	£2,925	£163	£73
29/01/2025	30	£150	17	56.7%	10	£225	8	80.0%	62.5%	£4,350	£174	£109
30/01/2025	30	£150	17	56.7%	10	£225	9	90.0%	65.0%	£4,575	£176	£114
31/01/2025	30	£150	29	96.7%	10	£225	10	100.0%	97.5%	£6,600	£169	£165

Figure 3.8: Month Ahead Dynamic Pricing - Using OCC% Forecast

3.4 Conclusions & Next Steps

3.4.1 Findings

Several attempts were made to create a SARIMA model, however processing capacity limitations prevented an accurate model from being generated. A simpler model was created to prove out the process of creation, evaluation and its use in one month ahead occupancy forecasting.

3.4.2 Next Steps

The next steps are primarily to refine the SARIMA and then explore the use of SARIMAX.

The main task in establishing an accurate SARIMA model is to obtain sufficient processing capacity in order to use AUTO ARIMA with a seasonal parameter of 365 and then refine this to achieve an accurate evaluation using the one year test data.

A SARIMAX model is then worth exploring to look at exogenous factors such as weather and economic indices as these are very likely to have an impact of hotel occupancy. Further refinements can also be extended to include tourism information such as flight searches.

4 Customer Satisfaction

4.1 Business Scope & Benefits

For the hotel operations within VGT, customer satisfaction is an important contributor to the overall efficiency and profitability of their business. This has been recognised by the hotel industry for several years: *“Service quality and customer satisfaction have gradually been recognized as key factors used to gain competitive advantage and customer retention”* (Yang, Jou and Cheng, 2011). If the levels of satisfaction across different aspects of a customers experience can be quantified, then action can be focused on improving those areas that are lacking. Clearly, identifying the areas that contribute most is part of the challenge. Not surprisingly, the use of data science to assist in this area is increasingly important (Zarezadeh, Rastegar and Xiang, 2022).

4.1.1 Measuring Customer Satisfaction

There are several aspects to quantifying customer satisfaction, for example:

- The specific areas of interest. Such as cleanliness, service experience, quality of facilities, local amenities and even things such as wifi availability
- Customer satisfaction ranking. Obtaining a customer’s views on how poor to excellent a specific area was
- Importance to the business. The appropriate weighting of the importance of each area.

It has become standard practice and easy to gather customer feedback on how they rate their experience after a stay at a hotel through check-out questionnaires and feedback requests (Li, Ye and Law, 2013), and also increasingly utilising reviews in services such as [Booking.Com](#) and [Tripadvisor](#). VGT already has some feedback mechanisms and some historical customer satisfaction data that can be used.

Previous studies have looked at a way of pulling the above aspects together and an example is the use of a matrix, an “I-S Model” (Yang, Jou and Cheng, 2011), see an example at Figure 4.1. This categorises the areas of interest and helps highlight those areas performing well and those not, whilst also reflecting their importance to the business. In particular, areas in the “To-be-improved” quadrant are important but have not met customer expectations and so these are areas that a hotel’s management team can focus on to improve the business.

4.2 Data Analysis Approach

In a similar way to Demand Forecasting, VGT does not currently formally analyse the experience of its customers across all hotels and then use this to inform management actions to improve its services. However, VGT does have a good history of customer ratings data that




Figure 4.1: Example of an Importance-Satisfaction Model

can be utilised initially to experiment with ways of more formally measuring customer satisfaction. If positive, then this can then be extended in sophistication and scope of data collection (see Section 4.4).

The objective of this initial experimentation step is to establish a process that can utilise existing historical customer review data for a single hotel and create an IS-Model to determine possible actions. In summary, the approach is to:

- Collate customer details and satisfaction ratings
- Utilise these to identify the areas that customers regard as important
- Establish a customer satisfaction prediction model
- Combine the above to create an IS-Model
- Use the IS-Model to identify actions that can improve customer service.

4.2.1 Data Required

The scope of the data is all the historical customer satisfaction ratings that VGT customers have provided after staying at a hotel. For consistency a Customer Satisfaction Score (CSAT) will be used, this is a 5 point Likert scale from “very dissatisfied” to “very satisfied” (Bishop, 2022). Data points required are:

- Overall customer rating
- Customer demographics (Age, Gender, Residence)
- Trip Background (Purpose, Travel Type, Booking Type)
- Individual service ratings (Amenities, Staff, Cleanliness, Wifi etc).

4.2.2 Classification Techniques to Determine Importance

The aim is to use a classification model to determine how the different factors contribute to the overall customer rating, ie to determine the relative importance of each factor. The target variable ‘y’ is the overall CSAT and the input feature set ‘X’ consists of the demographics, trip background and individual service ratings. The created model will provide an initial indication

of the relative importances across the feature set. The model can then also be used to provide predictions of the CSAT in response to value changes for individual features and this can help hotel management teams focus in planning what actions to take to improve customer service

Potential approaches for multi-class classification include:

- Support Vector Machine (SVM)
- K-Nearest Neighbour (KNN)
- Naive Bayes (NB)
- Decision Tree (DT)
- Random Forest (RF)
- Extreme Gradient Boosting (XGBoost).

A comparison of the classification of customer reviews, (Noori, 2021), examined several models, including: SVM, KNN, NB, DT; a DT approach was found to be the most accurate. Another examination of how to identify the important predictors of customer reviews found DT to be an accurate approach (Baouchi, 2018). Extending the DT approach, XGBoost was found to be a successful model in many situations (Chen and Guestrin, 2016). In this initial examination, XGBoost is used.

4.3 Simulated Data Analysis

4.3.1 Data Summary

A dataset was created to simulate the data as defined in the earlier section. The data elements are shown in the tables below and in more detail at Appendix C.

Table 4.1: Loaded File Data Types

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
id	103904	0	0	103904	int64	0	103904	0	0
gender	103904	0	0	2	category	103904	0	0	0
age	103904	0	0	75	int64	0	103904	0	0
purpose_of_travel	103904	0	0	5	category	103904	0	0	0
type_of_travel	103904	0	0	2	category	103904	0	0	0
type_of_booking	103904	0	0	3	category	103904	0	0	0
score_wifi	103904	0	0	4	int64	0	103904	0	0
score_transport	103904	0	0	5	int64	0	103904	0	0
score_booking	103904	0	0	5	int64	0	103904	0	0
score_location	103904	0	0	5	int64	0	103904	0	0
score_restaurant	103904	0	0	4	int64	0	103904	0	0
score_staff	103904	0	0	3	int64	0	103904	0	0
score_parking	103904	0	0	2	int64	0	103904	0	0
score_checkin	103904	0	0	5	int64	0	103904	0	0
score_local_sites	103904	0	0	5	int64	0	103904	0	0
score_housekeeping	103904	0	0	5	int64	0	103904	0	0
score_overall	103904	0	0	5	int64	0	103904	0	0

Table 4.2: Loaded File Characteristics

	count	mean	std	min	25%	50%	75%	max
id	103904.00	64924.21	37463.81	1.00	32533.75	64856.50	97368.25	129880.00
age	103904.00	39.38	15.11	7.00	27.00	40.00	51.00	85.00
score_wifi	103904.00	1.21	0.89	0.00	1.00	1.00	2.00	3.00
score_transport	103904.00	2.43	1.13	0.00	2.00	2.00	3.00	4.00
score_booking	103904.00	2.20	1.01	0.00	2.00	2.00	3.00	4.00
score_location	103904.00	2.33	0.91	0.00	2.00	2.00	3.00	4.00
score_restaurant	103904.00	1.54	0.96	0.00	1.00	1.00	2.00	3.00
score_staff	103904.00	1.14	0.59	0.00	1.00	1.00	2.00	2.00
score_parking	103904.00	0.24	0.43	0.00	0.00	0.00	0.00	1.00
score_checkin	103904.00	2.55	0.94	0.00	2.00	2.00	3.00	4.00
score_local_sites	103904.00	2.82	0.90	0.00	2.00	3.00	4.00	4.00
score_housekeeping	103904.00	2.57	0.97	0.00	2.00	2.00	3.00	4.00
score_overall	103904.00	2.21	1.30	0.00	1.00	2.00	3.00	4.00

4.3.2 XGBoost Model Creation & Evaluation

An XGBoost model was created using the above dataset randomly split 80:20 into training and testing data. The initial model was evaluated and a Confusion Matrix, Shap Values and F-Score were generated. The F-Score of 50% was not high which indicates the predictions are not very accurate and so more hyper-parameter tuning will be required; however, as an initial examination the model is still useful.

The relative feature importance ranking was calculated, shown in the figure below. Age clearly has a much larger importance than any of the other features.




Figure 4.2: XGBoost Feature Importance Results - All Features

It is useful to see what customer demographic factors and travel type impact the overall CSAT, however these cannot be directly changed by the hotel's management team, so a second model

was generated using just the individual service scores and this feature importance ranking is shown in the figure below. There is a closer spread of importance, with transport and booking being the most important.




Figure 4.3: XGBoost Feature Importance Results - Service Features Only

4.3.3 I-S Model

Using the predicted feature rankings, an I-S Model can be created. From this it is clear that several areas are in the “Excellent” quadrant and so these areas require no action, similarly parking is in the “Care-Free” quadrant and so it’s low satisfaction rating is not a priority to address. However, the service areas of restaurant and wifi are in the “To-be-Improved” quadrant which suggests that the hotel management team should investigate these to understand what the problems are and to see what improvements can be made.

	feature	importance	satisfaction_mean
0	score_transport	75	3.425912
1	score_booking	66	3.199935
2	score_location	65	3.333192
3	score_local_sites	63	3.818900
4	score_checkin	62	3.552443
5	score_wifi	61	2.209703
6	score_restaurant	57	2.539354
7	score_housekeeping	47	3.569901
8	score_staff	41	2.138532
9	score_parking	28	1.242657

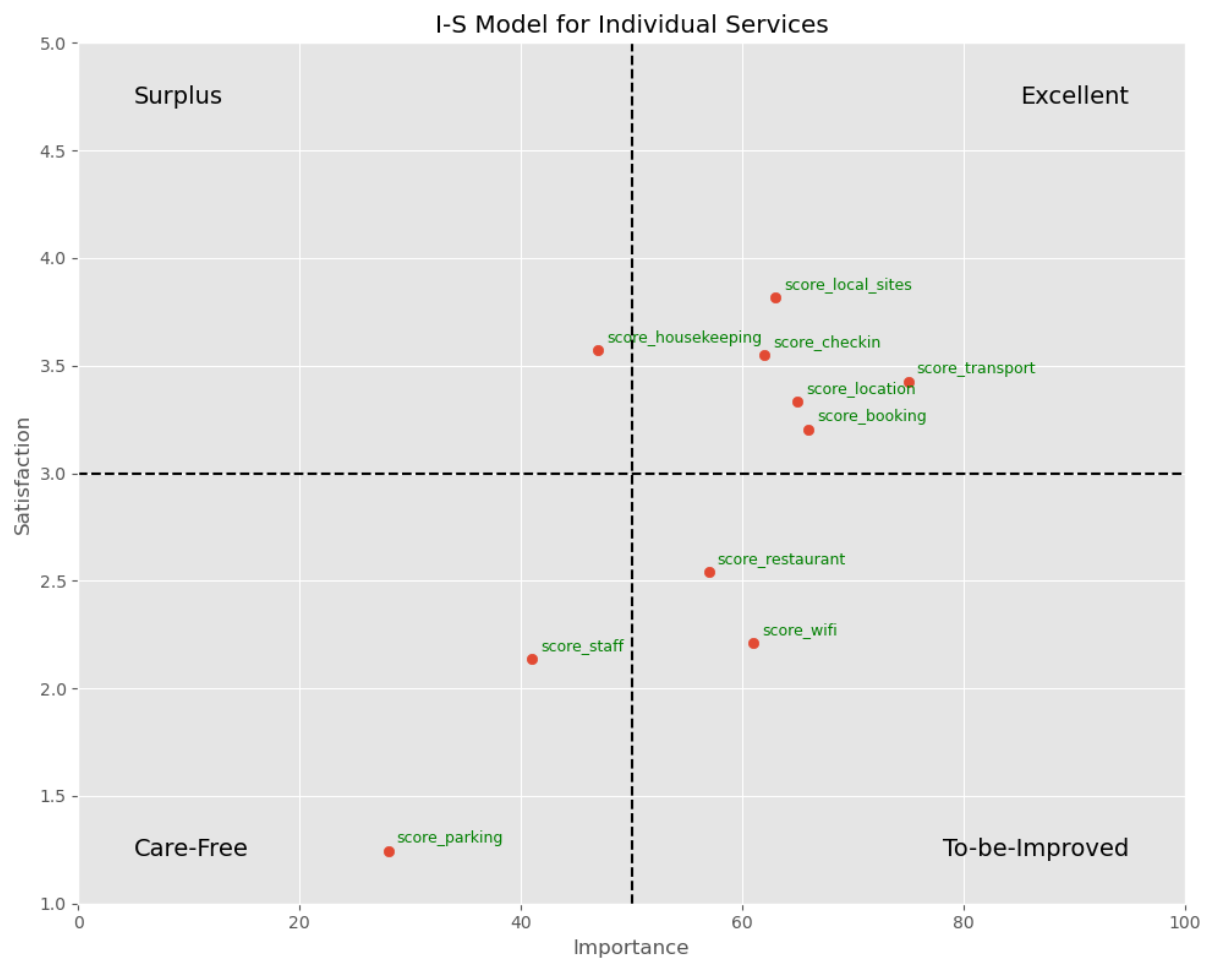


Figure 4.4: Importance-Significance Predictions

4.4 Conclusions & Next Steps

4.4.1 Findings

The use of the XGBoost model demonstrated that such a model is helpful in identifying which factors contribute most to overall customer satisfaction. However there does appear to be an over-correlation between service scores and importance once demographic factors are removed and further work on adjusting the feature set and model fine-tuning is needed.

4.4.2 Next Steps

The next steps can be considered in terms of: Model Refinement; Data Collection; Third-Party Data; and Business Processes.

The XGBoost model needs to be further refined to ensure that the importance rankings do reflect their real contribution. This requires a combination of: feature management as suggested earlier; and hyperparameter fine-tuning to improve accuracy. It may also be worthwhile using other techniques to cross-check the model results, for example: Principal Component Analysis or Logistic Regression.

The scope of features should be reviewed to ensure the appropriate areas are being considered. Internal systems and, for example, customer check-out processes can be changed to collect this data. Also sentiment analysis of written reviews can be used to derive an additional customer satisfaction metric.

Having utilised its own internal data on customer reviews, VGT will be like much of the hotel industry in not significantly incorporating third-party review data from sources such as [Tripadvisor](#) and a next-step is to investigate how to extract, analyse and incorporate this data (Park, 2023).

Once a model has been proven, this can be extended across all of VGT's hotels. The identification of action areas can then become part of standard management processes, with the result of the actions then tracked to see how customer-satisfaction improves over time.

References

- Andrew, W.P., Cranage, D.A. and Lee, C.K. (1990) 'Forecasting Hotel Occupancy Rates with Time Series Models: An Empirical Analysis', *Hospitality Research Journal*, 14(2), pp. 173–182. Available at: <https://doi.org/10.1177/109634809001400219>.
- Baouchi, S. (2018) *A text analytical approach: Predicting and understanding customer satisfaction by making use of customer reviews*. PhD thesis. Available at: <https://thesis.eur.nl/pub/53903/thesis-hotel-reviews-430727sb-laatsteversie-.pdf>.
- Bishop, C. (2022) 'What is customer satisfaction score? (+ how to measure CSAT)'. Available at: <https://www.zendesk.co.uk/blog/customer-satisfaction-score/>.
- Chen, T. and Guestrin, C. (2016) 'XGBoost: A scalable tree boosting system', in. New York, NY, USA: Association for Computing Machinery (Kdd '16), p. 785794. Available at: <https://doi.org/10.1145/2939672.2939785>.
- Chow, W.S., Shyu, J.-C. and Wang, K.-C. (1998) 'Developing a Forecast System for Hotel Occupancy Rate Using Integrated ARIMA Models', *Journal of International Hospitality, Leisure & Tourism Management*, 1(3), pp. 55–80. Available at: https://doi.org/10.1300/J268v01n03_05.
- FHA, F. (2023) 'Understanding occupancy rate in hotels: A comprehensive guide'. Available at: <https://fhahoreca.com/glossary/occupancy-rate/>.
- Huang, L. and Zheng, W. (2022) 'Hotel demand forecasting: a comprehensive literature review', *Tourism Review*, 78(1), pp. 218–244. Available at: <https://doi.org/10.1108/TR-07-2022-0367>.
- Jeffrey, D. and Hubbard, N.J. (1994) 'A model of hotel occupancy performance for monitoring and marketing in the hotel industry', *International Journal of Hospitality Management*, 13(1), pp. 57–71. Available at: [https://doi.org/10.1016/0278-4319\(94\)90059-0](https://doi.org/10.1016/0278-4319(94)90059-0).
- Johansson, A. (2022) 'Benefits and challenges in forecasting hotel demand'. Available at: <https://www.demandcalendar.com/blog/benefits-and-challenges-in-forecasting-hotel-demand>.
- Li, H., Ye, Q. and Law, R. (2013) 'Determinants of customer satisfaction in the hotel industry: An application of online review analysis', *Asia Pacific journal of tourism research*, 18(7), pp. 784–802. Available at: <https://www.tandfonline.com/doi/full/10.1080/10941665.2012.708351>.
- Lighthouse, A. (2024) 'Hotel ADR: Everything you need to know about Average Daily Rate'. Available at: <https://www.mylighthouse.com/resources/blog/hotel-adr-everything-you-need-to-know-about-average-daily-rate>.
- Lighthouse, B. (2023) 'The hotelier's ultimate guide to occupancy forecasting'. Available at:

<https://www.mylighthouse.com/resources/blog/how-to-forecast-hotel-occupancy>.

Noori, B. (2021) ‘Classification of customer reviews using machine learning algorithms’, *Applied Artificial Intelligence*, 35(8), pp. 567–588. Available at: <https://doi.org/10.1080/08839514.2021.1922843>.

Park, J. (2023) ‘Combined text-mining/DEA method for measuring level of customer satisfaction from online reviews’, *Expert Systems with Applications*, 232, p. 120767. Available at: <https://doi.org/10.1016/j.eswa.2023.120767>.

Phumchusri, N. and Suwatanapongched, P. (2023) ‘Forecasting hotel daily room demand with transformed data using time series methods’, *Journal of Revenue and Pricing Management*, 22(1), pp. 44–56. Available at: <https://doi.org/10.1057/s41272-021-00363-6>.

Weatherford, L.R. and Kimes, S.E. (2003) ‘A comparison of forecasting methods for hotel revenue management’, *International Journal of Forecasting*, 19(3), pp. 401–415. Available at: [https://doi.org/10.1016/S0169-2070\(02\)00011-0](https://doi.org/10.1016/S0169-2070(02)00011-0).

Yang, C.-C., Jou, Y.-T. and Cheng, L.-Y. (2011) ‘Using integrated quality assessment for hotel service quality’, *Quality & Quantity*, 45(2), pp. 349–364. Available at: <https://doi.org/10.1007/s11135-009-9301-4>.

Zarezadeh, Z.Z., Rastegar, R. and Xiang, Z. (2022) ‘Big data analytics and hotel guest experience: A critical analysis of the literature’, *International Journal of Contemporary Hospitality Management*, 34(6), pp. 2320–2336. Available at: <https://doi.org/10.1108/IJCHM-10-2021-1293>.

A Hotel Demand Forecasting - Jupyter Notebook Output

A.1 Data Load & Characteristics

- Load time series data and look at its characteristics
- Determine autocorrelation, seasonality, stationarity
- Decomposition,
- OLS Regression

A.1.1 Characteristics

Table A.1: Loaded File Data Types

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
Standard_OCC	1461	0	0	177	int64	0	1461	0	0
Standard_Capacity	1461	0	0	1	int64	0	1461	0	0
Standard_Rate	1461	0	0	1	int64	0	1461	0	0
Premium_OCC	1461	0	0	101	int64	0	1461	0	0
Premium_Capacity	1461	0	0	1	int64	0	1461	0	0
Premium_Rate	1461	0	0	1	int64	0	1461	0	0

	Standard_OCC	Standard_Capacity	Standard_Rate	Premium_OCC	\
Date					
2020-01-01	129	254	325	65	
2020-01-02	126	254	325	53	
2020-01-03	137	254	325	63	

	Premium_Capacity	Premium_Rate
Date		
2020-01-01	100	575
2020-01-02	100	575
2020-01-03	100	575

Table A.2: Loaded File Characteristics

	count	mean	std	min	25%	50%	75%	max
Standard_OCC	1461.00	135.65	50.22	38.00	88.00	136.00	183.00	254.00

	count	mean	std	min	25%	50%	75%	max
Standard_Capacity	1461.00	254.00	0.00	254.00	254.00	254.00	254.00	254.00
Standard_Rate	1461.00	325.00	0.00	325.00	325.00	325.00	325.00	325.00
Premium_OCC	1461.00	56.14	29.35	0.00	31.00	57.00	83.00	100.00
Premium_Capacity	1461.00	100.00	0.00	100.00	100.00	100.00	100.00	100.00
Premium_Rate	1461.00	575.00	0.00	575.00	575.00	575.00	575.00	575.00

A.1.2 Autocorrelation, Seasonality, Stationarity

- Determine autocorrelation, seasonality, stationarity .. lag plot, ACF plot, ADF test, Differencing
- Decomposition




Figure A.1: Occupancy for Each Room Category

- Shows definite annual seasonality with peak high and low seasons
- Also some infrequent spikes in bookings
- Possibly a small upward trend over time
- Premium rooms hit max and zero bookings several times ...
- Both categories of room show definite autocorrelation
- Premium rooms bunched up at max value and autocorrelation may be slightly less strong
- Some outliers when rooms are fully booked





Figure A.2: Lag Plots - Test




- Both exhibit strong autocorrelation that diminishes slowly after approximately 250 days
- A positive trend is suggested by the slowly diminishing autocorrelation
- Multiple peaks at 350 days indicates annual seasonality
- ?? carry out PACF to determine the autoregressive order does indicate that it is autoregressive

ADF Test for Standard Rooms

ADF Statistic: -2.362457191578427
p-value: 0.15261408046089647
Critical Value 1%: -3.434908816804013
Critical Value 5%: -2.863553406963303
Critical Value 10%: -2.5678419239852994
Conclusion: Non-Stationary

ADF Test for Premium Rooms
ADF Statistic: -2.1359470749871265
p-value: 0.2303078418058474
Critical Value 1%: -3.434911997169608
Critical Value 5%: -2.863554810504947
Critical Value 10%: -2.567842671398422
Conclusion: Non-Stationary

Differenced Time Series



ADF Test for Standard Rooms
ADF Statistic: -4.650429185341465
p-value: 0.00010417359492157454
Critical Value 1%: -3.4349151819757466
Critical Value 5%: -2.863556216004778
Critical Value 10%: -2.5678434198545568
Conclusion: Stationary

ADF Test for Premium Rooms
ADF Statistic: -6.091009300062941
p-value: 1.0365104637060615e-07
Critical Value 1%: -3.4349151819757466
Critical Value 5%: -2.863556216004778
Critical Value 10%: -2.5678434198545568

Conclusion: Stationary

- Confirms that both time series are non-stationary




Figure A.3: Occupancy Time Series Decomposition

- Both have a small positive trend with room occupancy increasing 5 to 105% pa
- Confirms both time series are seasonal, with annual peaks and troughs
- On top of the annual seasonality, there are regular spikes leading to 100% occupancy
- Close clustering of residuals with some outliers that correspond to the seasonal spikes

A.2 Ordinary Least Squares (OLS) Linear Regression

- Unlikely to be a good model for forecasting given the strong seasonality, but examine to confirm

Durbin-Watson statistic: 0.07265854622866856

Durbin-Watson statistic: 0.24391749000237792

- The two fitted lines do not capture any seasonality
- Also show a downward trend line, which is not consistent with the decomposition trend line
- The Durban Watson statistics for both time series are less than 1.5 which confirms evidence of autocorrelation




Figure A.4: Occupancy Time Series Decomposition




Figure A.5: OLS Linear Regression

A.3 SARIMA Model Creation & Forecasting

- The strong autocorrelation for the room occupancy time series, suggest an autoregressive model such as ARMA or ARIMA or SARIMA
- The occupancy history is non-stationary and there is an upward trend, so the integrated component of ARIMA will automatically perform the differencing needed to transform the data into stationary data. So preferable to ARMA
- The strong seasonality suggests SARIMA would be most appropriate as it can process seasonal patterns. So preferable to ARIMA

SARIMA

- nb to be clear SARIMA stands for Seasonal Autoregressive Integrated Moving Average
- S: seasonal component, here handle the annual occupancy seasonality
- AR: autoregressive component p
- I: integrated, here handle the positive historical trend through differencing to make it stationary d
- MA: moving average component q
- SARIMA(p,d,q)(P,D,Q),m
- p,d,q non-seasonal
- P,D,Q seasonal
- m the seasonality period, length of the seasonal cycle
- !! But this is daily, with an annual cycle so would be m=365 which is too large ??

Residuals etc

- Residuals plot
- ACF plot of residuals
- Durbin-Watson of residuals
- Use accuracy measures MAE, RMSE, MAPE

Approach

- Confirm/Assess the autoregressive order (AR, p) and moving average order (MA, q). Using PACF plot and ACF plot respectively [repeat autocorrelation findings from previously?]
- Use the SARIMA(p,d,q)(P,D,Q,m) model
- Identify the best parameters using Auto Arima and using AIC (Akaike Information Criterion) to compare

A.3.1 Data Load & SARIMA Model Factors

- Load time series data
- Assess the autoregressive order (AR, p) - Using PACF plot
- Assess moving average order (MA, q) - Using ACF plot

PACF plot suggests AR order, p of approximately 3


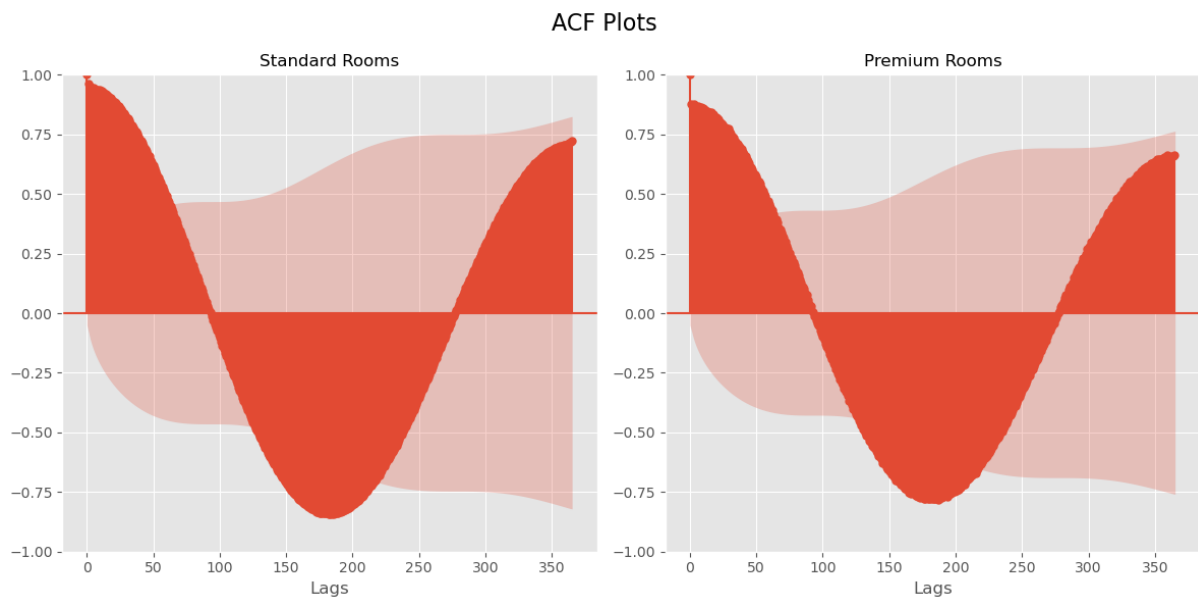


Figure A.6: Partial Autocorrelation Function Plot



ACF plot suggests MA order q of approximately 60

A.3.2 Create the SARIMA Model

- Save the training dataset to be used to train a SARIMA model, in a separate notebook
- Load the model in here for forecasting

A.3.3 Forecast Using the SARIMA Model




Figure A.7: One Month Occupancy Forecast

B SARIMA Models Creation + Auto ARIMA - Jupyter Notebook Output

B.1 SARIMA Model Create & Evaluate - Manual

Create the fitted model and save in model_sarima_manual

SARIMA Model Manual Creation using order: (3, 1, 1) and seasonal order: (3, 1, 1, 7)

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 9 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.08611D+00 |proj g|= 3.32757D-01

At iterate 5 f= 3.96176D+00 |proj g|= 5.15657D-02

At iterate 10 f= 3.89847D+00 |proj g|= 2.28767D-02

At iterate 15 f= 3.87075D+00 |proj g|= 2.68053D-02

At iterate 20 f= 3.84927D+00 |proj g|= 3.93821D-03

At iterate 25 f= 3.84921D+00 |proj g|= 2.54060D-03

At iterate 30 f= 3.84892D+00 |proj g|= 1.21032D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
9	34	38	1	0	0	1.154D-05	3.849D+00

F = 3.8489137053819098

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Residuals Analysis for Model:

SARIMAX Results

```
=====
Dep. Variable:          Standard_OCC      No. Observations:          1096
Model:                SARIMAX(3, 1, 1)x(3, 1, 1, 7)      Log Likelihood          -4218.409
Date:                  Sun, 12 Jan 2025      AIC                  8454.819
Time:                  15:59:43      BIC                  8499.748
Sample:                01-01-2020      HQIC                 8471.825
                        - 12-31-2022
=====
```


Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.1063      0.033      3.209      0.001      0.041      0.171
ar.L2          0.0429      0.027      1.582      0.114     -0.010      0.096
ar.L3         -0.0970      0.023     -4.245      0.000     -0.142     -0.052
ma.L1         -0.7384      0.031    -23.567      0.000     -0.800     -0.677
ar.S.L7        -0.1167      0.035     -3.297      0.001     -0.186     -0.047
ar.S.L14       -0.0196      0.045     -0.436      0.663     -0.108      0.069
ar.S.L21       -0.0310      0.042     -0.746      0.456     -0.113      0.051
ma.S.L7        -0.8523      0.030    -28.508      0.000     -0.911     -0.794
sigma2        135.0472      2.162     62.464      0.000    130.810    139.285
=====
```


```
=====
Ljung-Box (L1) (Q):          0.02      Jarque-Bera (JB):          22690.36
Prob(Q):                    0.88      Prob(JB):              0.00
Heteroskedasticity (H):      1.67      Skew:                  2.07
Prob(H) (two-sided):        0.00      Kurtosis:              24.98
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



Model Evaluation on Testing Data



RMSE is: 260.18

MAPE is: 258.63

Stored 'model_sarima_manual' (SARIMAXResultsWrapper)

B.2 SARIMA Model Create & Evaluate - Auto ARIMA - First

Save in model_sarima_auto

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,1,1)[7]	: AIC=inf, Time=4.31 sec
ARIMA(0,1,0)(0,1,0)[7]	: AIC=9473.205, Time=0.03 sec
ARIMA(1,1,0)(1,1,0)[7]	: AIC=8920.643, Time=0.13 sec
ARIMA(0,1,1)(0,1,1)[7]	: AIC=inf, Time=0.63 sec
ARIMA(1,1,0)(0,1,0)[7]	: AIC=9250.699, Time=0.04 sec
ARIMA(1,1,0)(2,1,0)[7]	: AIC=8802.692, Time=0.56 sec
ARIMA(1,1,0)(2,1,1)[7]	: AIC=inf, Time=2.03 sec
ARIMA(1,1,0)(1,1,1)[7]	: AIC=inf, Time=0.68 sec
ARIMA(0,1,0)(2,1,0)[7]	: AIC=9022.330, Time=0.10 sec
ARIMA(2,1,0)(2,1,0)[7]	: AIC=8765.968, Time=0.70 sec
ARIMA(2,1,0)(1,1,0)[7]	: AIC=8880.910, Time=0.34 sec
ARIMA(2,1,0)(2,1,1)[7]	: AIC=inf, Time=2.03 sec
ARIMA(2,1,0)(1,1,1)[7]	: AIC=inf, Time=1.57 sec
ARIMA(3,1,0)(2,1,0)[7]	: AIC=8730.320, Time=0.81 sec
ARIMA(3,1,0)(1,1,0)[7]	: AIC=8844.538, Time=0.55 sec
ARIMA(3,1,0)(2,1,1)[7]	: AIC=inf, Time=3.39 sec
ARIMA(3,1,0)(1,1,1)[7]	: AIC=inf, Time=1.89 sec
ARIMA(4,1,0)(2,1,0)[7]	: AIC=8693.506, Time=0.98 sec
ARIMA(4,1,0)(1,1,0)[7]	: AIC=8807.355, Time=0.47 sec
ARIMA(4,1,0)(2,1,1)[7]	: AIC=inf, Time=4.74 sec
ARIMA(4,1,0)(1,1,1)[7]	: AIC=inf, Time=2.94 sec
ARIMA(5,1,0)(2,1,0)[7]	: AIC=8676.272, Time=1.08 sec
ARIMA(5,1,0)(1,1,0)[7]	: AIC=8794.927, Time=0.73 sec
ARIMA(5,1,0)(2,1,1)[7]	: AIC=inf, Time=4.92 sec
ARIMA(5,1,0)(1,1,1)[7]	: AIC=inf, Time=2.64 sec
ARIMA(5,1,1)(2,1,0)[7]	: AIC=8616.898, Time=2.21 sec
ARIMA(5,1,1)(1,1,0)[7]	: AIC=8718.876, Time=1.47 sec
ARIMA(5,1,1)(2,1,1)[7]	: AIC=inf, Time=6.63 sec
ARIMA(5,1,1)(1,1,1)[7]	: AIC=8451.943, Time=4.63 sec
ARIMA(5,1,1)(0,1,1)[7]	: AIC=inf, Time=2.34 sec
ARIMA(5,1,1)(1,1,2)[7]	: AIC=8453.944, Time=6.92 sec
ARIMA(5,1,1)(0,1,0)[7]	: AIC=9024.409, Time=0.54 sec
ARIMA(5,1,1)(0,1,2)[7]	: AIC=8452.035, Time=3.76 sec
ARIMA(5,1,1)(2,1,2)[7]	: AIC=8455.726, Time=11.32 sec
ARIMA(4,1,1)(1,1,1)[7]	: AIC=8449.945, Time=3.87 sec
ARIMA(4,1,1)(0,1,1)[7]	: AIC=inf, Time=2.36 sec
ARIMA(4,1,1)(1,1,0)[7]	: AIC=8717.998, Time=1.41 sec
ARIMA(4,1,1)(2,1,1)[7]	: AIC=inf, Time=6.14 sec
ARIMA(4,1,1)(1,1,2)[7]	: AIC=8451.941, Time=6.10 sec
ARIMA(4,1,1)(0,1,0)[7]	: AIC=9024.170, Time=0.55 sec
ARIMA(4,1,1)(0,1,2)[7]	: AIC=8450.036, Time=3.55 sec
ARIMA(4,1,1)(2,1,0)[7]	: AIC=8615.480, Time=2.36 sec
ARIMA(4,1,1)(2,1,2)[7]	: AIC=8453.741, Time=8.85 sec
ARIMA(3,1,1)(1,1,1)[7]	: AIC=8451.302, Time=2.96 sec
ARIMA(4,1,2)(1,1,1)[7]	: AIC=8446.786, Time=4.72 sec
ARIMA(4,1,2)(0,1,1)[7]	: AIC=inf, Time=4.19 sec
ARIMA(4,1,2)(1,1,0)[7]	: AIC=8718.912, Time=2.10 sec
ARIMA(4,1,2)(2,1,1)[7]	: AIC=inf, Time=10.23 sec
ARIMA(4,1,2)(1,1,2)[7]	: AIC=8444.931, Time=7.78 sec

```

ARIMA(4,1,2)(0,1,2)[7] : AIC=8446.244, Time=6.86 sec
ARIMA(4,1,2)(2,1,2)[7] : AIC=8454.879, Time=9.97 sec
ARIMA(3,1,2)(1,1,2)[7] : AIC=8519.228, Time=8.83 sec
ARIMA(5,1,2)(1,1,2)[7] : AIC=8455.934, Time=7.30 sec
ARIMA(4,1,3)(1,1,2)[7] : AIC=inf, Time=10.00 sec
ARIMA(3,1,1)(1,1,2)[7] : AIC=8453.293, Time=6.55 sec
ARIMA(3,1,3)(1,1,2)[7] : AIC=8456.475, Time=9.54 sec
ARIMA(5,1,3)(1,1,2)[7] : AIC=inf, Time=10.30 sec
ARIMA(4,1,2)(1,1,2)[7] intercept : AIC=8453.667, Time=8.94 sec

```

Best model: ARIMA(4,1,2)(1,1,2)[7]

Total fit time: 223.717 seconds

SARIMA Model Manual Creation using order: (4, 1, 2) and seasonal order: (1, 1, 2, 7)

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 10 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.03105D+00 |proj g|= 3.18759D-01

At iterate 5 f= 3.91558D+00 |proj g|= 2.75819D-02

At iterate 10 f= 3.85716D+00 |proj g|= 8.66397D-02

At iterate 15 f= 3.84864D+00 |proj g|= 1.49384D-03

At iterate 20 f= 3.84831D+00 |proj g|= 1.25557D-03

At iterate 25 f= 3.84801D+00 |proj g|= 8.50538D-03

At iterate 30 f= 3.84767D+00 |proj g|= 1.84933D-03

At iterate 35 f= 3.84759D+00 |proj g|= 2.58246D-03

At iterate 40 f= 3.84749D+00 |proj g|= 3.01463D-03

At iterate 45 f= 3.84706D+00 |proj g|= 2.46234D-02

At iterate 50 f= 3.84349D+00 |proj g|= 1.32273D-02

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
10	50	59	1	0	0	1.323D-02	3.843D+00

F = 3.8434904043823033

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT

Residuals Analysis for Model:

SARIMAX Results

Dep. Variable:	Standard_OCC	No. Observations:	1096
Model:	SARIMAX(4, 1, 2)x(1, 1, 2, 7)	Log Likelihood	-4212.465
Date:	Sun, 12 Jan 2025	AIC	8444.931
Time:	16:20:40	BIC	8494.852
Sample:	01-01-2020	HQIC	8463.827
	- 12-31-2022		


Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	1.1291	0.039	28.783	0.000	1.052	1.206
ar.L2	-0.0889	0.026	-3.358	0.001	-0.141	-0.037
ar.L3	-0.1674	0.032	-5.157	0.000	-0.231	-0.104
ar.L4	0.0423	0.025	1.707	0.088	-0.006	0.091
ma.L1	-1.8247	0.038	-48.012	0.000	-1.899	-1.750
ma.L2	0.8513	0.031	27.169	0.000	0.790	0.913
ar.S.L7	0.7630	0.217	3.511	0.000	0.337	1.189
ma.S.L7	-1.6872	0.209	-8.079	0.000	-2.097	-1.278
ma.S.L14	0.7191	0.183	3.932	0.000	0.361	1.078
sigma2	132.6832	1.995	66.524	0.000	128.774	136.592


Ljung-Box (L1) (Q):	1.87	Jarque-Bera (JB):	25082.63
Prob(Q):	0.17	Prob(JB):	0.00
Heteroskedasticity (H):	1.76	Skew:	2.29
Prob(H) (two-sided):	0.00	Kurtosis:	26.07

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



Model Evaluation on Testing Data



RMSE is: 311.34

MAPE is: 307.80

Stored 'model_sarima_auto' (SARIMAXResultsWrapper)

'\n#\n# Convert the pmdarima model to a statsmodels SARIMAX model for residual analysis\nmod

B.3 SARIMA Model Create & Evaluate - Auto ARIMA - Second

Save in model_sarima_auto_second

Performing stepwise search to minimize aic

```

ARIMA(0,1,0)(0,1,0)[7]      : AIC=9473.205, Time=0.04 sec
ARIMA(1,1,0)(1,1,0)[7]      : AIC=8920.643, Time=0.15 sec
ARIMA(0,1,1)(0,1,1)[7]      : AIC=inf, Time=0.48 sec
ARIMA(1,1,0)(0,1,0)[7]      : AIC=9250.699, Time=0.08 sec
ARIMA(1,1,0)(2,1,0)[7]      : AIC=8802.692, Time=0.32 sec
ARIMA(1,1,0)(3,1,0)[7]      : AIC=8723.919, Time=0.53 sec
ARIMA(1,1,0)(3,1,1)[7]      : AIC=inf, Time=2.71 sec
ARIMA(1,1,0)(2,1,1)[7]      : AIC=inf, Time=2.17 sec
ARIMA(0,1,0)(3,1,0)[7]      : AIC=8941.142, Time=0.53 sec
ARIMA(2,1,0)(3,1,0)[7]      : AIC=8694.720, Time=0.74 sec
ARIMA(2,1,0)(2,1,0)[7]      : AIC=8765.968, Time=0.42 sec
ARIMA(2,1,0)(3,1,1)[7]      : AIC=inf, Time=4.20 sec
ARIMA(2,1,0)(2,1,1)[7]      : AIC=inf, Time=1.94 sec
ARIMA(3,1,0)(3,1,0)[7]      : AIC=8654.898, Time=1.04 sec
ARIMA(3,1,0)(2,1,0)[7]      : AIC=8730.320, Time=0.66 sec
ARIMA(3,1,0)(3,1,1)[7]      : AIC=inf, Time=4.86 sec
ARIMA(3,1,0)(2,1,1)[7]      : AIC=inf, Time=3.19 sec
ARIMA(3,1,1)(3,1,0)[7]      : AIC=8541.718, Time=2.65 sec
ARIMA(3,1,1)(2,1,0)[7]      : AIC=8613.852, Time=1.79 sec
ARIMA(3,1,1)(3,1,1)[7]      : AIC=8454.819, Time=5.42 sec
ARIMA(3,1,1)(2,1,1)[7]      : AIC=inf, Time=3.95 sec
ARIMA(3,1,1)(3,1,2)[7]      : AIC=inf, Time=7.51 sec
ARIMA(3,1,1)(2,1,2)[7]      : AIC=8455.068, Time=7.00 sec
ARIMA(2,1,1)(3,1,1)[7]      : AIC=8459.822, Time=4.41 sec
ARIMA(3,1,2)(3,1,1)[7]      : AIC=inf, Time=10.65 sec
ARIMA(2,1,2)(3,1,1)[7]      : AIC=inf, Time=10.53 sec
ARIMA(3,1,1)(3,1,1)[7] intercept : AIC=8456.706, Time=12.51 sec

```

Best model: ARIMA(3,1,1)(3,1,1)[7]

Total fit time: 90.579 seconds

SARIMAX Results

```

=====
Dep. Variable:          y      No. Observations:      1096
Model:          SARIMAX(3, 1, 1)x(3, 1, 1, 7)      Log Likelihood      -4218.409
Date:          Wed, 08 Jan 2025      AIC      8454.819
Time:          17:54:27      BIC      8499.748
Sample:          01-01-2020      HQIC      8471.825
              - 12-31-2022

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.1063      0.033      3.209      0.001      0.041      0.171
ar.L2          0.0429      0.027      1.582      0.114     -0.010      0.096

```


ar.L3	-0.0970	0.023	-4.245	0.000	-0.142	-0.052
ma.L1	-0.7384	0.031	-23.567	0.000	-0.800	-0.677
ar.S.L7	-0.1167	0.035	-3.297	0.001	-0.186	-0.047
ar.S.L14	-0.0196	0.045	-0.436	0.663	-0.108	0.069
ar.S.L21	-0.0310	0.042	-0.746	0.456	-0.113	0.051
ma.S.L7	-0.8523	0.030	-28.508	0.000	-0.911	-0.794
sigma2	135.0472	2.162	62.464	0.000	130.810	139.285

=====

Ljung-Box (L1) (Q):	0.02	Jarque-Bera (JB):	22690.36
Prob(Q):	0.88	Prob(JB):	0.00
Heteroskedasticity (H):	1.67	Skew:	2.07
Prob(H) (two-sided):	0.00	Kurtosis:	24.98

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
 SARIMA Model Manual Creation using order: (3, 1, 1) and seasonal order: (3, 1, 1, 7)
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 9 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.08611D+00 |proj g|= 3.32757D-01

At iterate 5 f= 3.96176D+00 |proj g|= 5.15657D-02

At iterate 10 f= 3.89847D+00 |proj g|= 2.28767D-02

At iterate 15 f= 3.87075D+00 |proj g|= 2.68053D-02

At iterate 20 f= 3.84927D+00 |proj g|= 3.93821D-03

At iterate 25 f= 3.84921D+00 |proj g|= 2.54060D-03

At iterate 30 f= 3.84892D+00 |proj g|= 1.21032D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
9	34	38	1	0	0	1.154D-05	3.849D+00

F = 3.8489137053819098

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Residuals Analysis for Model:

SARIMAX Results

```
=====
Dep. Variable:                Standard_OCC      No. Observations:      1096
Model:                SARIMAX(3, 1, 1)x(3, 1, 1, 7)      Log Likelihood      -4218.409
Date:                Wed, 08 Jan 2025      AIC      8454.819
Time:                17:54:34      BIC      8499.748
Sample:                01-01-2020      HQIC      8471.825
                        - 12-31-2022
=====
```


Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.1063      0.033      3.209      0.001      0.041      0.171
ar.L2          0.0429      0.027      1.582      0.114     -0.010      0.096
ar.L3         -0.0970      0.023     -4.245      0.000     -0.142     -0.052
ma.L1         -0.7384      0.031    -23.567      0.000     -0.800     -0.677
ar.S.L7        -0.1167      0.035     -3.297      0.001     -0.186     -0.047
ar.S.L14       -0.0196      0.045     -0.436      0.663     -0.108      0.069
ar.S.L21       -0.0310      0.042     -0.746      0.456     -0.113      0.051
ma.S.L7        -0.8523      0.030    -28.508      0.000     -0.911     -0.794
sigma2        135.0472      2.162     62.464      0.000    130.810    139.285
=====
```


```
=====
Ljung-Box (L1) (Q):                0.02      Jarque-Bera (JB):                22690.36
Prob(Q):                0.88      Prob(JB):                0.00
Heteroskedasticity (H):            1.67      Skew:                2.07
Prob(H) (two-sided):            0.00      Kurtosis:            24.98
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



Model Evaluation on Testing Data



RMSE is: 260.18

MAPE is: 258.63

Stored 'model_sarima_auto_second' (SARIMAXResultsWrapper)

C Customer Satisfaction - Jupyter Notebook Output

C.1 Data Load & Characteristics

- Load data and look at its characteristics

C.1.1 Libraries & Functions

C.1.2 Characteristics

Table C.1: Loaded File Data Types

	Count	Missing	Empty	Unique	Type	String	Int	Float	List
id	103904	0	0	103904	int64	0	103904	0	0
gender	103904	0	0	2	category	103904	0	0	0
age	103904	0	0	75	int64	0	103904	0	0
purpose_of_travel	103904	0	0	5	category	103904	0	0	0
type_of_travel	103904	0	0	2	category	103904	0	0	0
type_of_booking	103904	0	0	3	category	103904	0	0	0
score_wifi	103904	0	0	4	int64	0	103904	0	0
score_transport	103904	0	0	5	int64	0	103904	0	0
score_booking	103904	0	0	5	int64	0	103904	0	0
score_location	103904	0	0	5	int64	0	103904	0	0
score_restaurant	103904	0	0	4	int64	0	103904	0	0
score_staff	103904	0	0	3	int64	0	103904	0	0
score_parking	103904	0	0	2	int64	0	103904	0	0
score_checkin	103904	0	0	5	int64	0	103904	0	0
score_local_sites	103904	0	0	5	int64	0	103904	0	0
score_housekeeping	103904	0	0	5	int64	0	103904	0	0
score_overall	103904	0	0	5	int64	0	103904	0	0

```

      id  gender  age  purpose_of_travel  type_of_travel  type_of_booking  \
0   70172   Male   13          aviation  Personal Travel      Not defined
1    5047   Male   25          tourism    Group Travel    Group bookings
2  110028  Female   26          tourism    Group Travel    Group bookings

```

```

      score_wifi  score_transport  score_booking  score_location  \
0              1                3                2                1

```

1	1	2	2	2
2	1	2	2	2
	score_restaurant	score_staff	score_parking	score_checkin \
0	3	2	1	3
1	0	0	0	1
2	3	2	1	3
	score_local_sites	score_housekeeping	score_overall	
0	4	4	2	
1	3	1	0	
2	3	4	4	

Table C.2: Loaded File Characteristics

	count	mean	std	min	25%	50%	75%	max
id	103904.00	64924.21	37463.81	1.00	32533.75	64856.50	97368.25	129880.00
age	103904.00	39.38	15.11	7.00	27.00	40.00	51.00	85.00
score_wifi	103904.00	1.21	0.89	0.00	1.00	1.00	2.00	3.00
score_transport	103904.00	2.43	1.13	0.00	2.00	2.00	3.00	4.00
score_booking	103904.00	2.20	1.01	0.00	2.00	2.00	3.00	4.00
score_location	103904.00	2.33	0.91	0.00	2.00	2.00	3.00	4.00
score_restaurant	103904.00	1.54	0.96	0.00	1.00	1.00	2.00	3.00
score_staff	103904.00	1.14	0.59	0.00	1.00	1.00	2.00	2.00
score_parking	103904.00	0.24	0.43	0.00	0.00	0.00	0.00	1.00
score_checkin	103904.00	2.55	0.94	0.00	2.00	2.00	3.00	4.00
score_local_sites	103904.00	2.82	0.90	0.00	2.00	3.00	4.00	4.00
score_housekeeping	103904.00	2.57	0.97	0.00	2.00	2.00	3.00	4.00
score_overall	103904.00	2.21	1.30	0.00	1.00	2.00	3.00	4.00

C.2 XGBoost Model Creation & Forecasting


C.2.1 Initial XGBoost Model Creation


```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=True, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, objective='multi:softprob', ...)
```

C.2.2 Evaluate The XGBoost Model

	precision	recall	f1-score	support
0	0.94	0.98	0.96	2142
1	0.47	0.47	0.47	4826
2	0.45	0.46	0.45	4745
3	0.47	0.48	0.48	4589
4	0.46	0.43	0.44	4479
accuracy			0.51	20781
macro avg	0.56	0.56	0.56	20781
weighted avg	0.51	0.51	0.51	20781

Classification Error: 0.49





C.2.3 XGBoost Model - Features Ranking





Figure C.1: XGBoost Feature Importance Results - All Features

C.2.4 XGBoost Model - Features Subset

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=True, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, objective='multi:softprob', ...)
```

	precision	recall	f1-score	support
0	0.93	0.98	0.95	2142
1	0.46	0.44	0.45	4826
2	0.44	0.49	0.46	4745
3	0.46	0.44	0.45	4589
4	0.44	0.41	0.43	4479
accuracy			0.50	20781
macro avg	0.54	0.55	0.55	20781
weighted avg	0.50	0.50	0.50	20781

Classification Error: 0.50






Figure C.2: XGBoost Feature Importance Results - Service Features Only

	feature	importance	satisfaction_mean
0	score_transport	75	3.425912
1	score_booking	66	3.199935

	feature	importance	satisfaction_mean
2	score_location	65	3.333192
3	score_local_sites	63	3.818900
4	score_checkin	62	3.552443
5	score_wifi	61	2.209703
6	score_restaurant	57	2.539354
7	score_housekeeping	47	3.569901
8	score_staff	41	2.138532
9	score_parking	28	1.242657

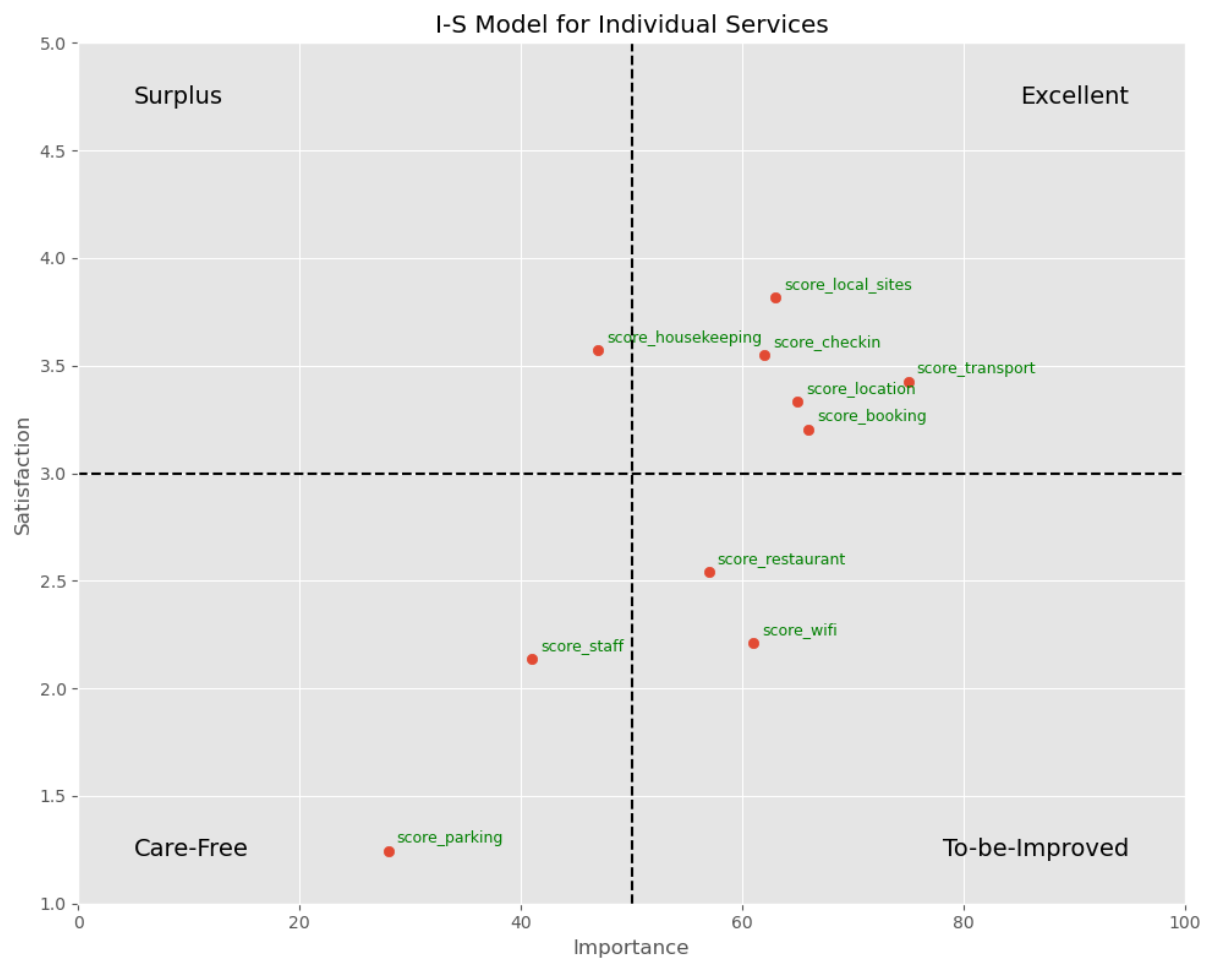


Figure C.3: Importance-Significance Predictions